

Northwind Home: operating the agent, and the fire drill

Domain FDE Pre-Work Deep Track • Blue Data Consulting • Stage 7: Operate

Before you start

Read [usecase1_day7_problem_statement.pdf](#) first: it sets out a week in production, and why the agent can change with no code change at all. This brief is the practical: how the agent is released, watched and rolled back in Google ADK, and how you will prove you can do it under pressure.

The technical question: **when nothing in the code changed, what are you watching?** Today you release the agent as one versioned unit, watch its behaviour and its service levels, and then someone breaks it.

Today's decision

What the agent must keep doing, how you will know within minutes that it stopped, and exactly how you get back.

The release: everything that shapes behaviour

Part	In the kit	Why it is versioned
Model	release.yaml: model	A new model changes behaviour
Instructions	releases/r1/instructions/	One removed sentence can change a promise
Policies	releases/r1/policies/	What the knowledge agent believes is true
Tool settings	release.yaml: tool_settings	A tool can be switched off, or fail
Fingerprints	release.yaml: fingerprints	Any edit made outside a release shows as drift

The agent in production

Part	What it does
v7_operations	The Day 5 coordinator, with the Day 6 grounded assistant as a fourth specialist: knowledge_agent
OpsPlugin	Stamps the release on every trace span, and records every tool failure
Retry plugin	When a tool fails, lets the model retry twice, then report the failure instead of crashing
Trace tab	In adk web: every model call and tool call, with its timing and its release

How you watch it

Signal	Target (SLO)	Where it comes from
Must-never contract cases	All 3 pass. No budget	ops/contract.yaml, run by ops.check
Should-usually contract cases	At least 4 of 5	The same run
Tool error rate	At most 5%	tool_error events in runs.jsonl
Answer time, 95th percentile	Within 30 seconds	turn events in runs.jsonl
Tokens per model call	At most 6,000	model_call events in runs.jsonl

What was built

Nothing the agent does is new today. What is new is everything around it.

	Day 6	Day 7
Agents	Two teams: one acts, one answers	One production app: the team, with the knowledge agent inside it
Change	Edit a file	Create a release, pass the check, promote it, and roll back if needed
Evidence	An eval run	Traces stamped with the release, a health check, and a release history
Failure	A failed test	A tool that fails mid-conversation, handled and counted against a budget
Proof	A before and after score	A fire drill: detected, rolled back, recovered, and written up

Your task: 50 minutes, in pairs

Pairs. Work in usecase01_day7. Every command is in README-training-account.md. Record everything in Day7_Worksheets_UC01.docx. During the drill, one of you breaks the agent and the other fixes it; then swap.

#	Part	Time	What you do
1	Predict it cold	10 min	For each of the four kinds of change, predict which signal would show it first.
2	Look inside	10 min	Run the tests, read the release, make and undo a drift, read a trace, run the health check.
3	Fire drill	20 min	Your partner deploys a hidden fault. Detect it, diagnose it, roll back, confirm recovery. Swap.
4	Prove it	10 min	Write the incident notes, the one-page plan, and architecture v2.

Seven questions

1. Which faults you met did the behaviour contract catch, and which showed only in the SLIs, the traces or the release diff?
2. Did any fault pass every contract case? If so, why, and what does that say about graceful degradation?
3. How long did it take from deploy to detect, and from detect to recovered? What would make both shorter?
4. A model update arrives under the same name. How would you even know, and what would you watch?
5. Who may create a release, who may promote it, and who may roll it back without asking?
6. Which SLO would you tighten first, and which has no error budget at all?
7. A customer complains about a promise made last Tuesday. How do you find exactly what the agent was told that day?

Bring to the debrief

Your fastest detection, and the signal that gave it away · the fault that nearly got past you · one change you would make to the SLOs.

All data is made up. Nothing changes in any real system. If the environment breaks, hand up.